

# Data Analysis and Machine Learning using Python

Lecture 6: Boosted Decision Trees and Multi-Layer Perceptrons;  
intro to scikit-learn  
*April 20 2024*

# Today:

- Quick review of mid-term quiz
- Reminder: Basic ideas of supervised machine learning
- Boosted Decision Trees
- Perceptrons and Neural Networks
- How to use scikit-learn

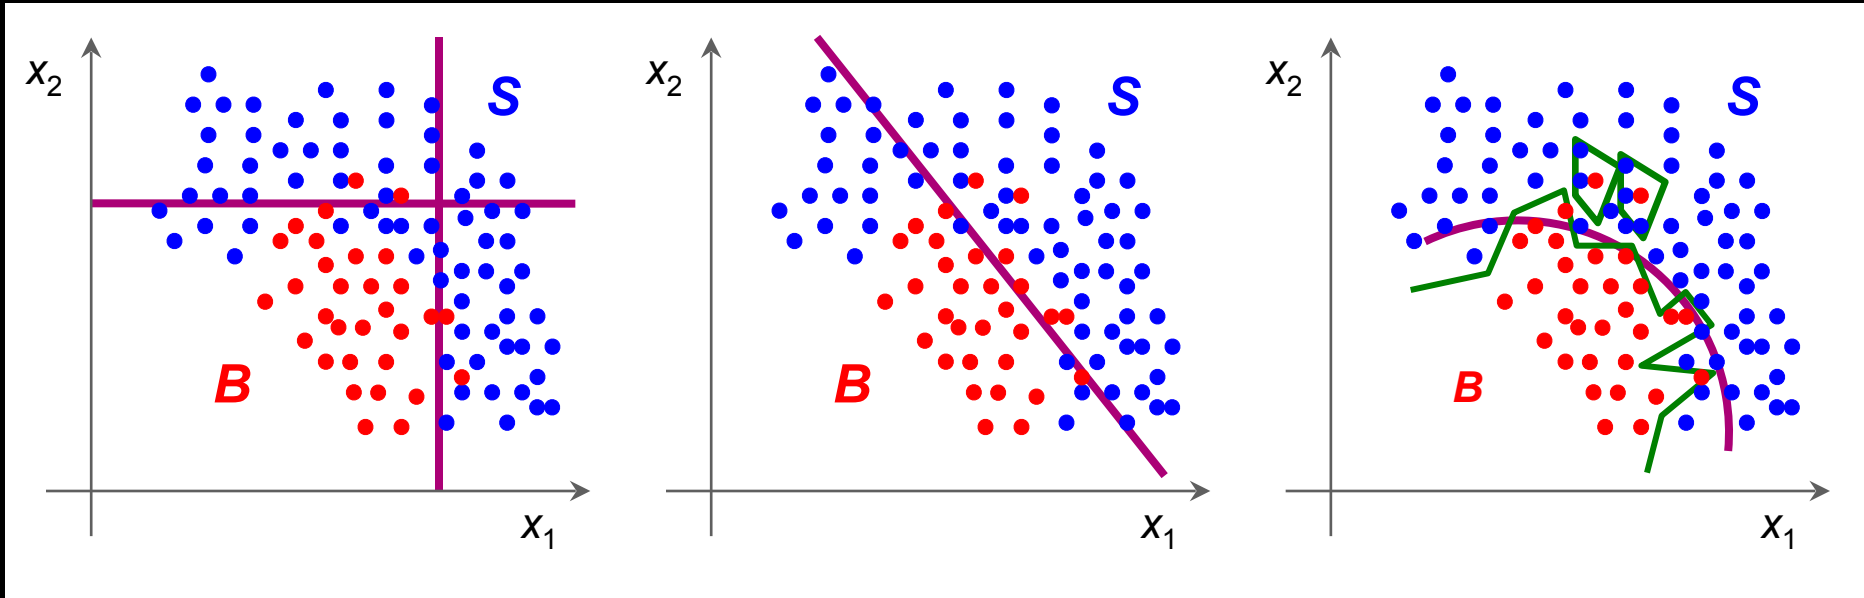
# Multivariate Analysis: Classification and Regression

- Related, common problems in data analysis
- **Multivariate analysis:** Multiple features (input variables)
- **Classification:** Which *class* does a particular object belong to, based on features?
- **Regression:** Predict output for a set of features (e.g., interpolation/extrapolation)

# Supervised machine learning

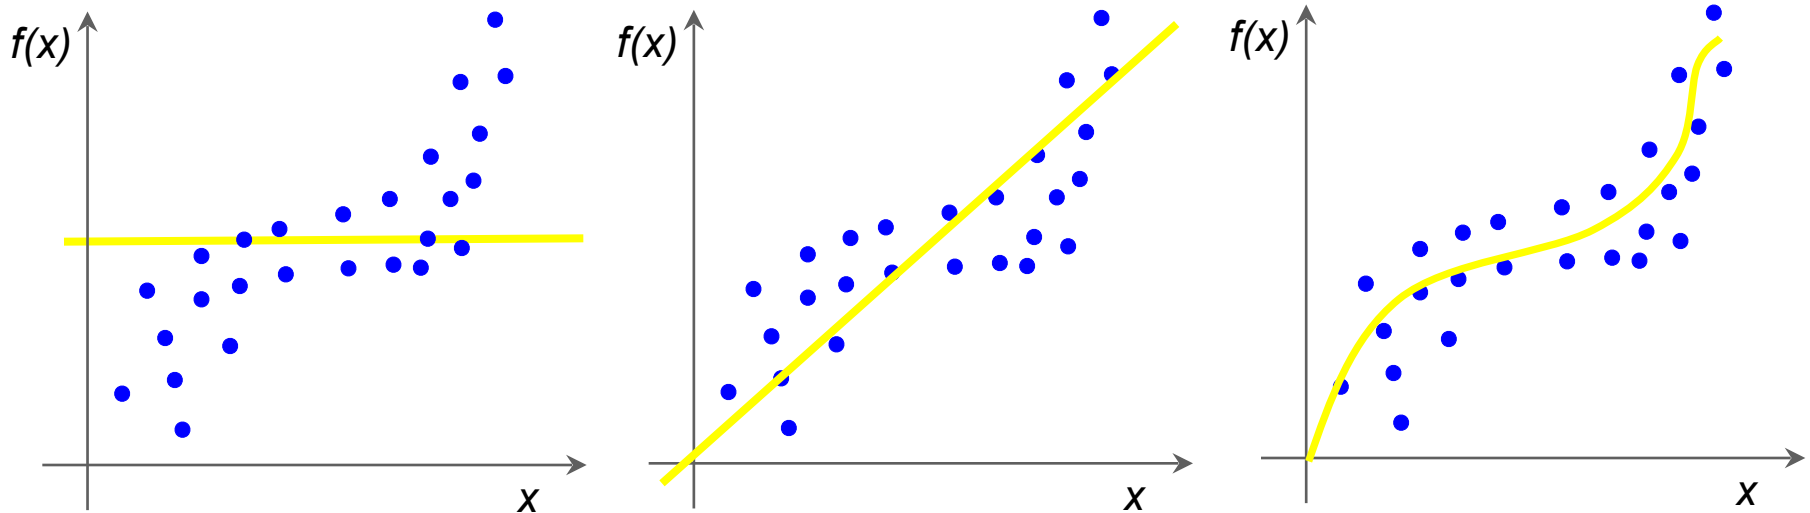
- **Machine learning (ML):** Build *model* for classification or regression automatically
- **Supervised ML:** Model “learns” from training sample for which correct answer is known
- **Model:** “function” determined by ML algorithm and used for classification or regression
  - Model provides output value for any valid input
  - Model has to generalize from the training sample to test sample

# Classification



- Classification: Find “decision boundary” to separate different event classes (in physics, often “signal” and “background”)
- Boundary:  $D-1$  dimensional hypersurface in  $D$ -dimensional feature space

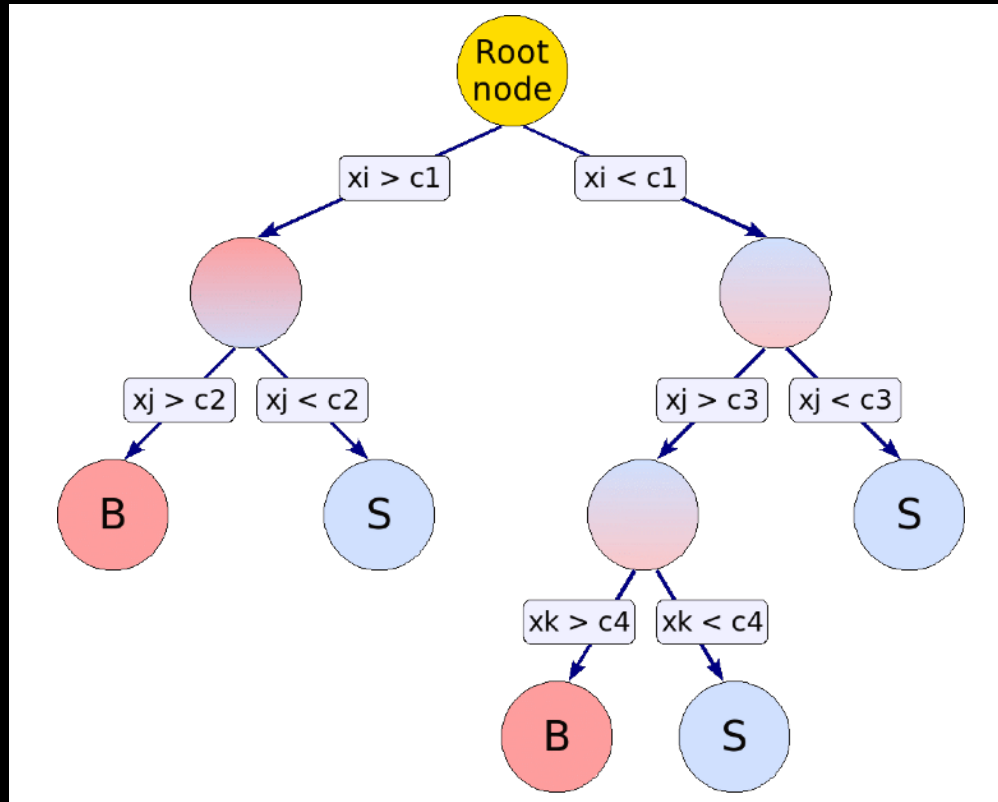
# Regression



- Regression: Find “model” to predict  $f(\mathbf{x})$  for every point  $\mathbf{x}$  in feature space
  - in general,  $\mathbf{x}$  is  $D$ -dimensional vector
- Model:  $D$ -dimensional “hypersurface” in  $(D+1)$ -dimensional space  $(\mathbf{x}, f(\mathbf{x}))$

# Boosted Decision Trees

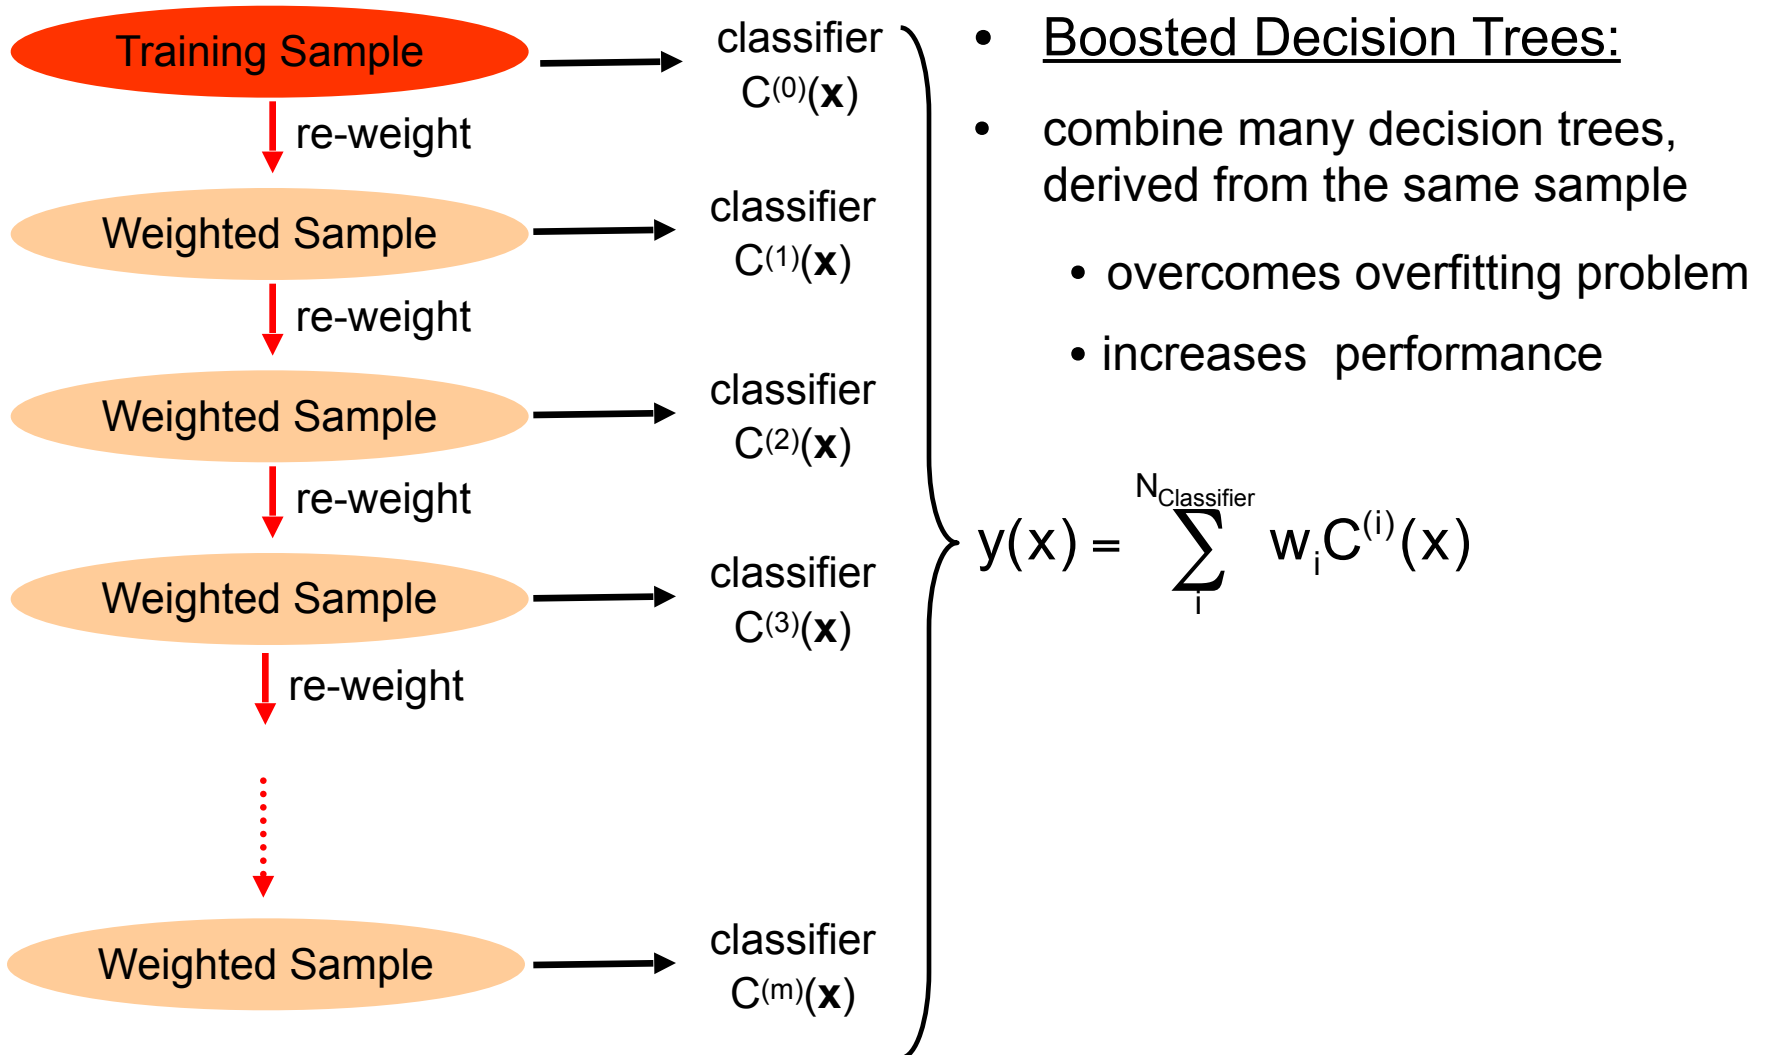
# Decision Trees



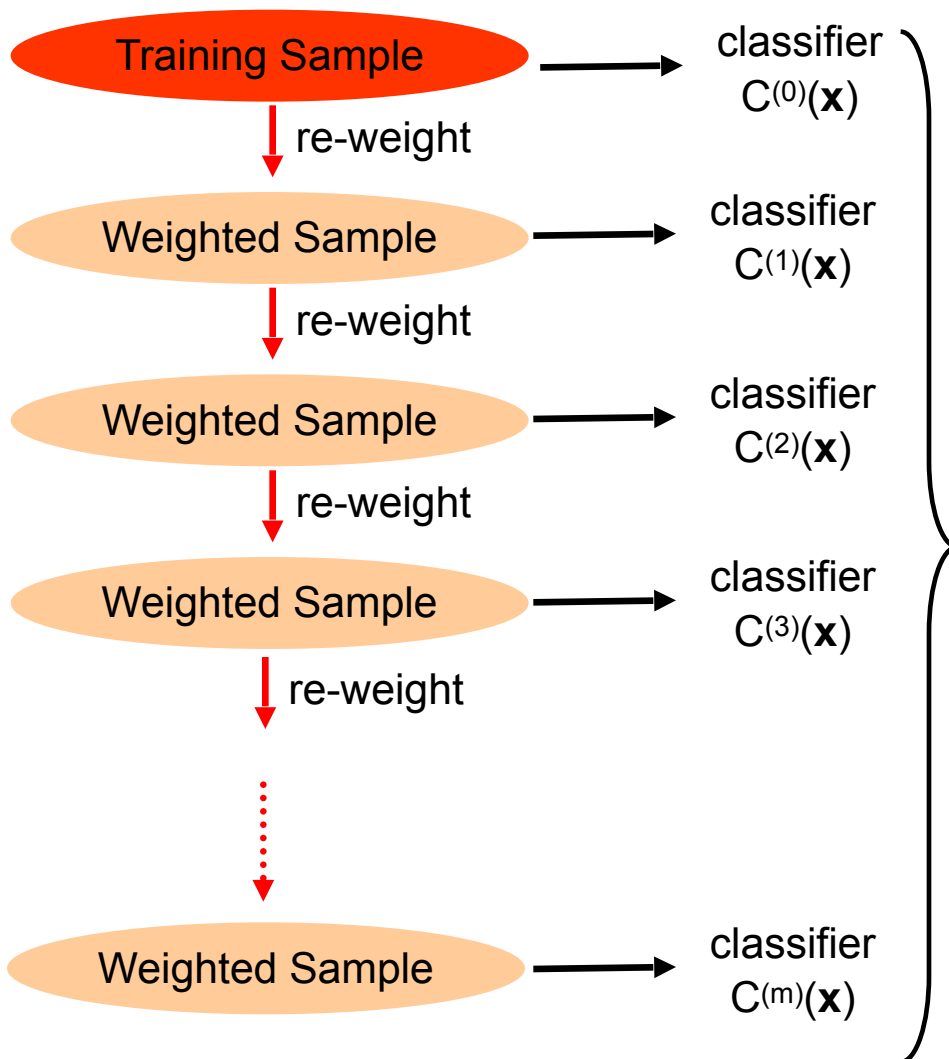
- Sequential application of cuts
- Each cut on different variable
- Final cut in each branch splits data into signal and background



# Boosting



# Adaptive Boosting (AdaBoost)



- AdaBoost re-weights events misclassified by previous classifier by:

$$\frac{1 - f_{\text{err}}}{f_{\text{err}}} \quad \text{with :}$$

$$f_{\text{err}} = \frac{\text{misclassified events}}{\text{all events}}$$

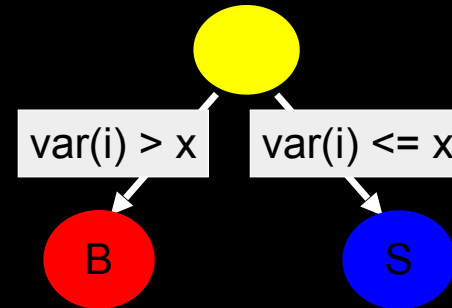
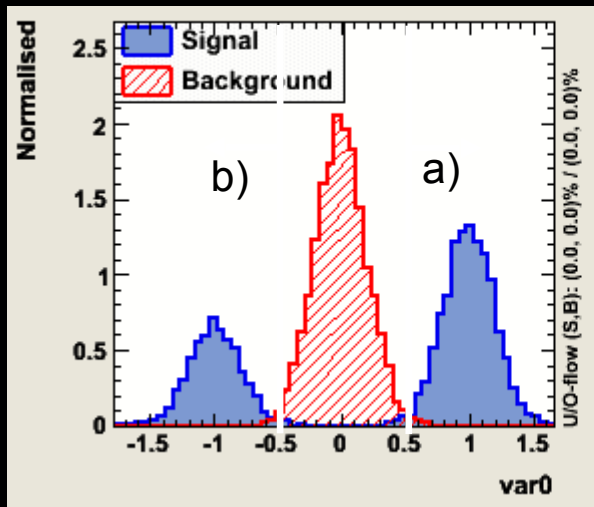
- AdaBoost weights the classifiers also using the error rate of the individual classifier according to:

$$y(\mathbf{x}) = \sum_i^{N_{\text{Classifier}}} \log\left(\frac{1 - f_{\text{err}}^{(i)}}{f_{\text{err}}^{(i)}}\right) C^{(i)}(\mathbf{x})$$

# AdaBoost: A simple demonstration

Example:

- Data file with three “bumps”
- Weak classifier (i.e. one single simple “cut”  $\leftrightarrow$  decision tree stumps )



a)  $\text{Var0} > 0.5 \rightarrow \epsilon_{\text{sig}}=66\% \epsilon_{\text{bkg}} \approx 0\%$  misclassified events in total 16.5%

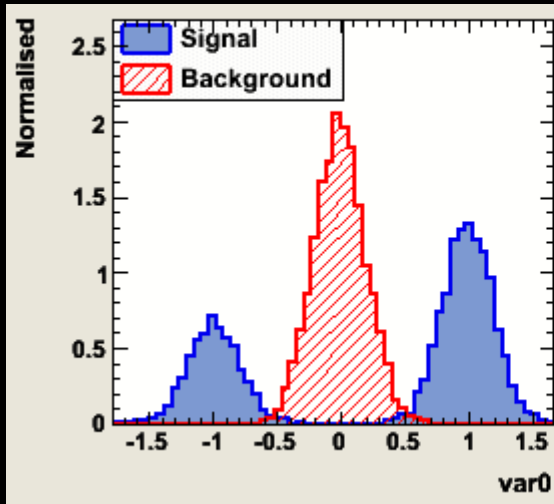
or

b)  $\text{Var0} < -0.5 \rightarrow \epsilon_{\text{sig}}=33\% \epsilon_{\text{bkg}} \approx 0\%$  misclassified events in total 33%

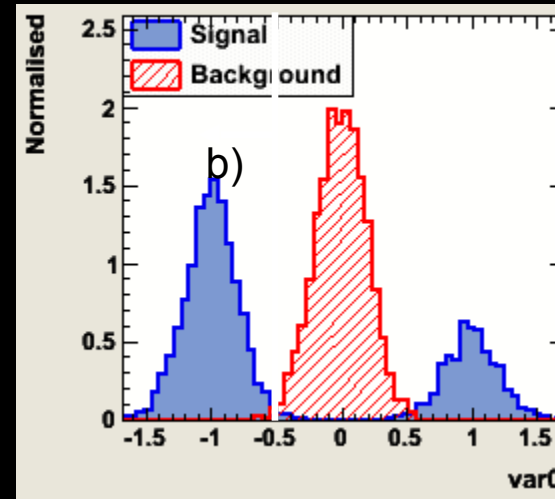
the training of a single decision tree stump will find “cut a)”

# AdaBoost: A simple demonstration

- before building the next “tree”: weight wrong classified training events by  $(1 - \text{err}/\text{err}_{\text{best}}) \approx 5$
- the next “tree” sees essentially the following data sample:

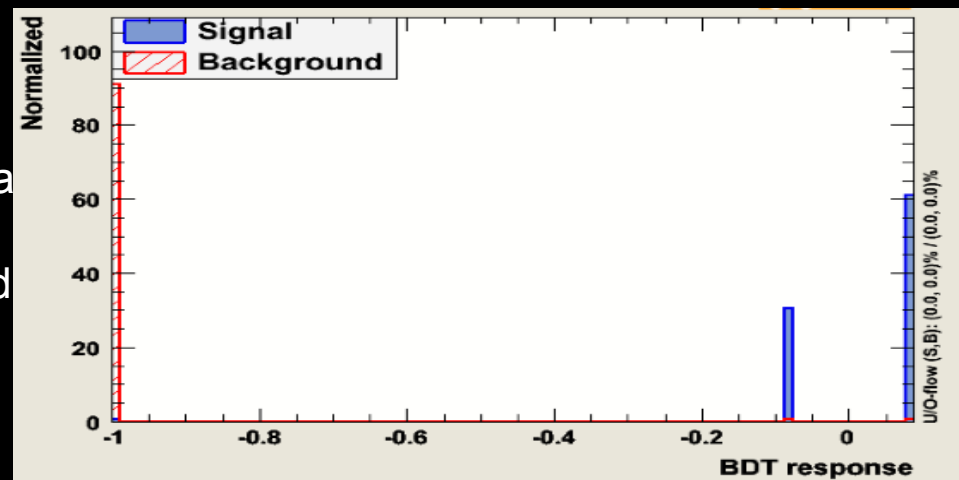


re-weight

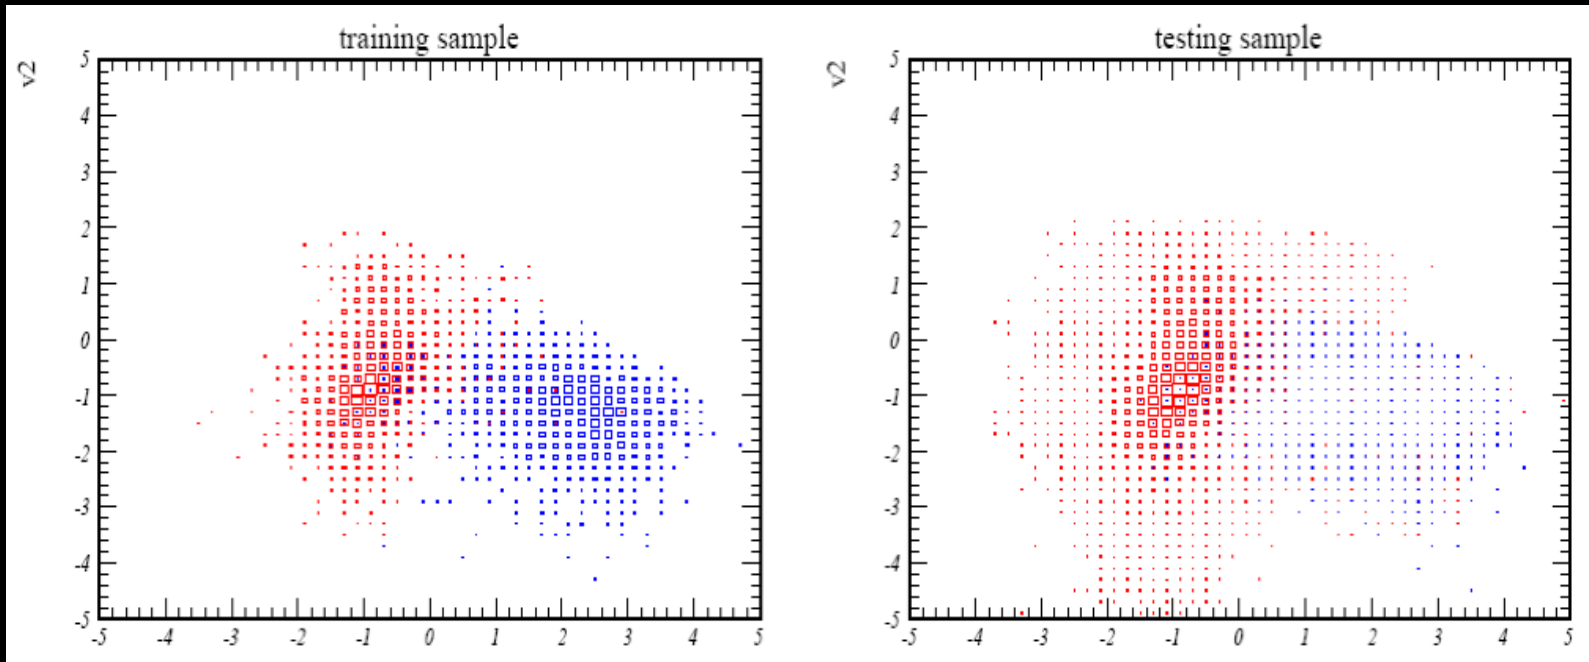


.. and hence will  
chose: “cut b)”:  
 $\text{Var0} < -0.5$

The combined classifier: Tree1 + Tree2  
the (weighted) average of the response to a  
test event from both trees is able to  
separate signal from background as good  
as one would expect from the most  
powerful classifier



# Training and test sample



- training and test sample do not have to be identical
- Can train on different ratios of signal and background events

# General Advice for (MVA) Analyses

There is no magic in MVA-Methods:

- you typically still need to make careful tuning and do some “hard work”
- no “artificial intelligence” ... just “fitting decision boundaries” in a given model

The most important thing at the start is finding good observables

- good separation power between S and B

- little correlations amongst each other

- no correlation with the parameters you try to measure in your signal sample!

Think also about possible combination of variables

- this may allow you to eliminate correlations

- rem.: you are MUCH more intelligent than your computer

Apply pure preselection cuts and let the MVA only do the difficult part.

“Sharp features should be avoided” → numerical problems, loss of information when binning is applied

- simple variable transformations (i.e.  $\log(\text{variable})$ ) can often smooth out these areas and allow signal and background differences to appear in a clearer way

Treat regions in the detector that have different features “independent”

- can introduce correlations where otherwise the variables would be uncorrelated!

Let's look at some BDT examples in  
Scikit learn

# Artificial Neural Networks



# Artificial Neural Networks

## Human brain

- Neuron switching time  $\sim .001$  second
- Number of neurons  $\sim 10^{10}$
- Connections per neuron  $\sim 10^{4-5}$
- Scene recognition time  $\sim .1$  second
- Requires a lot of training
- Lots of parallel computation!

## Artificial neural networks (ANNs):

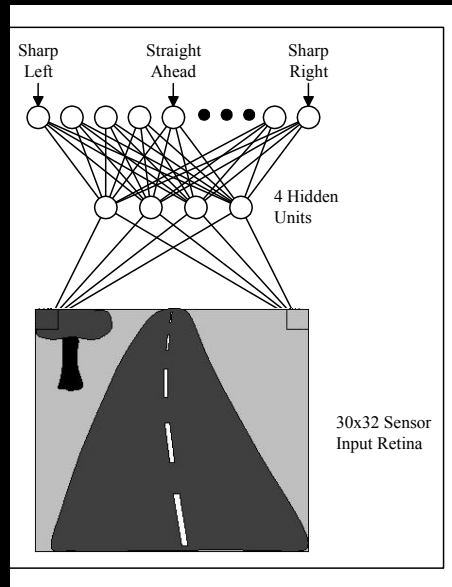
- Many neuron-like threshold switching units

# When to Consider Neural Networks

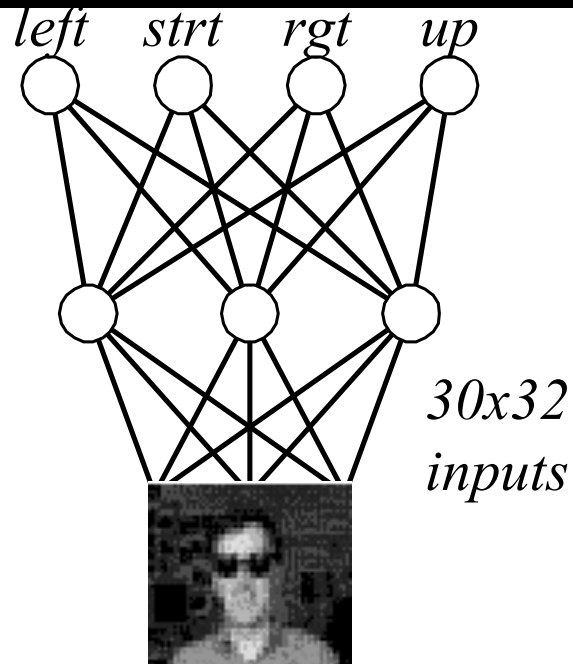
- Input is high-dimensional discrete or real-valued (e.g., combining information from different measurements or detectors)
- Output is discrete or real valued
- Output is a vector of values
- Possibly noisy data
- Form of target function is unknown
- Human readability of result is *unimportant*

Examples:

# ANN drives through Paris

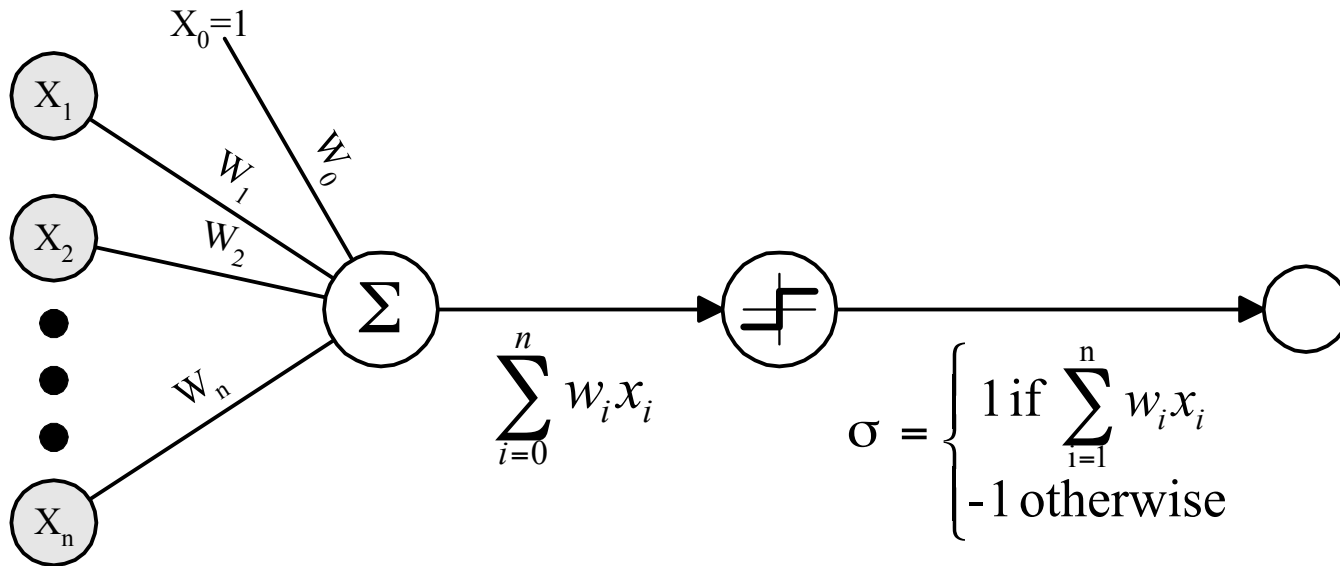


# ANN recognizes faces



*Typical Input Images*

# Basic unit: Perceptron

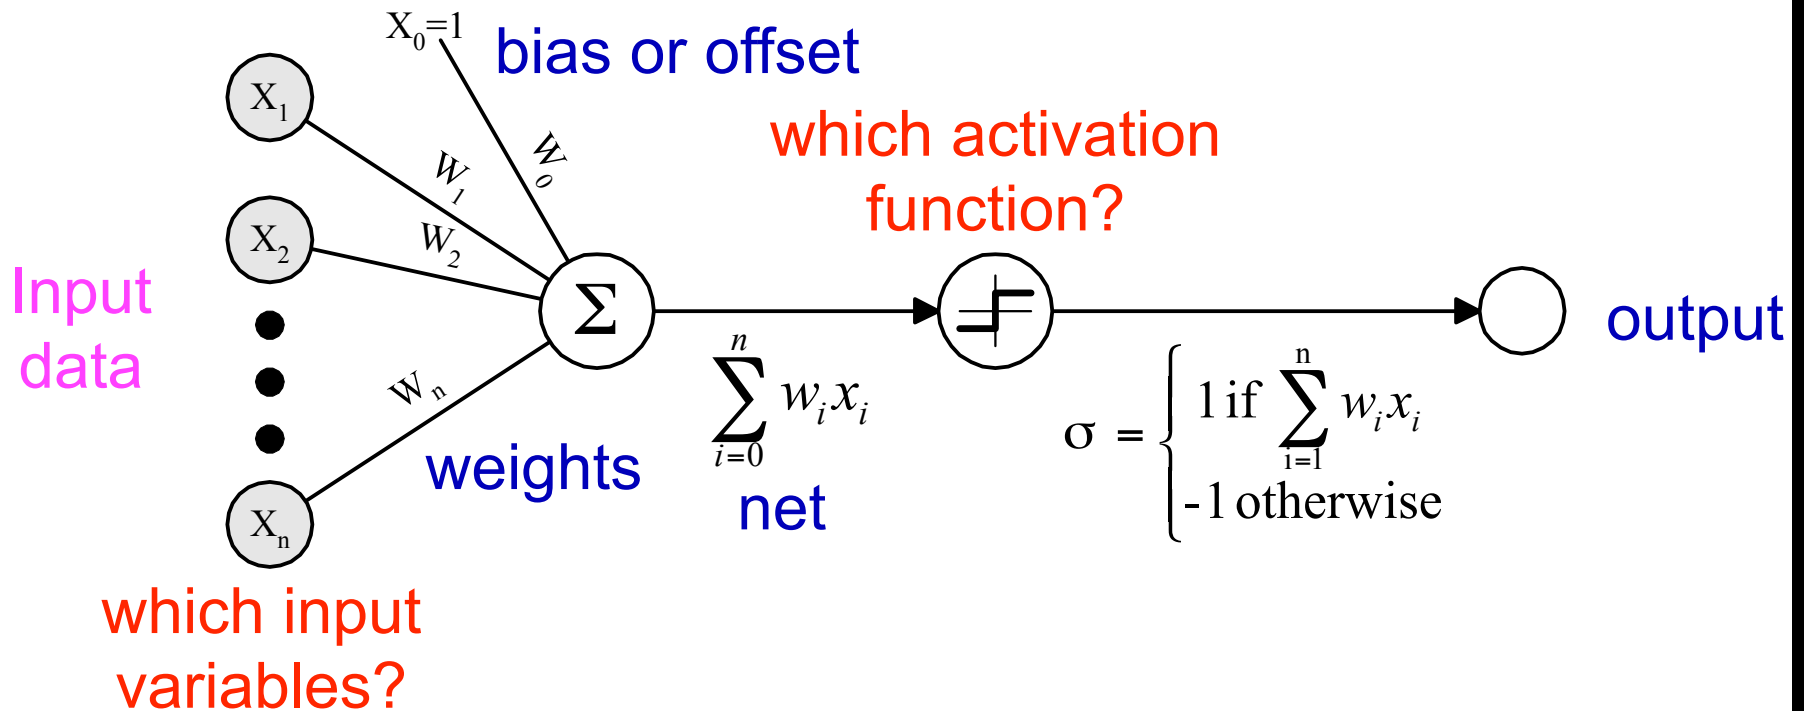


$$o(x_1, \dots, x_n) = \begin{cases} 1 & \text{if } w_0 + w_1 x_1 + \dots + w_n x_n > 0 \\ -1 & \text{otherwise} \end{cases}$$

Sometimes we will use simpler vector notation :

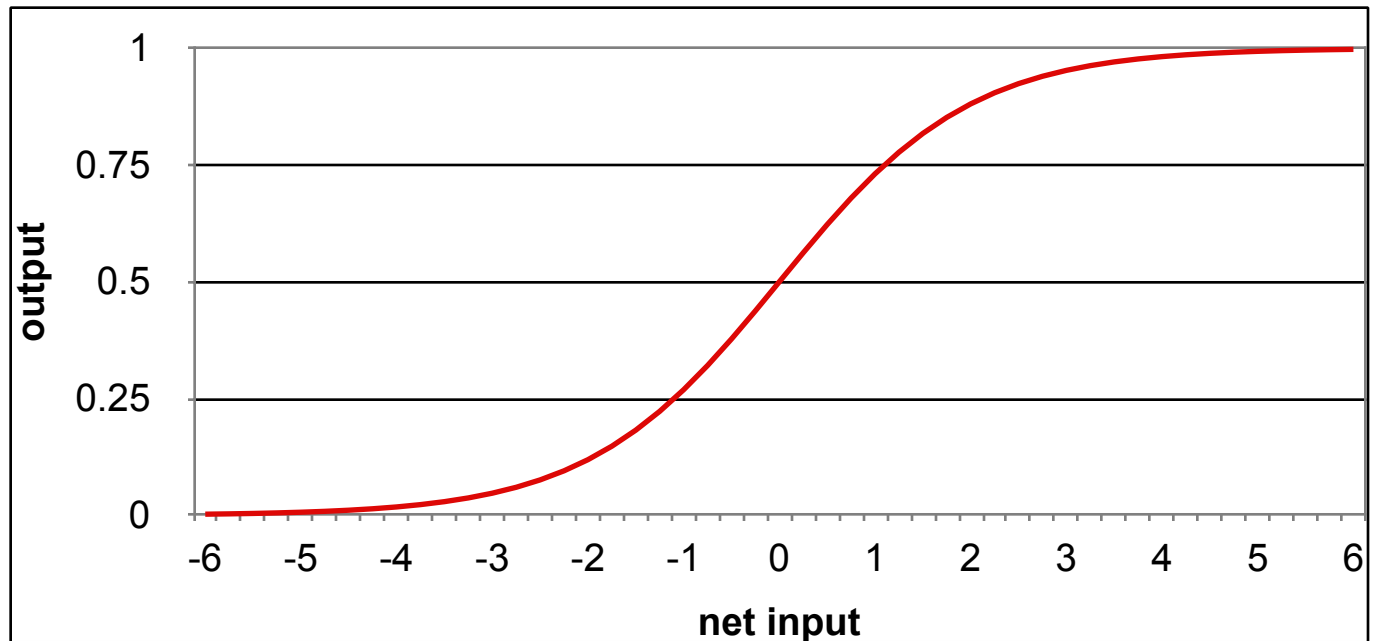
$$o(\vec{x}) = \begin{cases} 1 & \text{if } \vec{w} \cdot \vec{x} > 0 \\ -1 & \text{otherwise} \end{cases}$$

# Basic unit: Perceptron



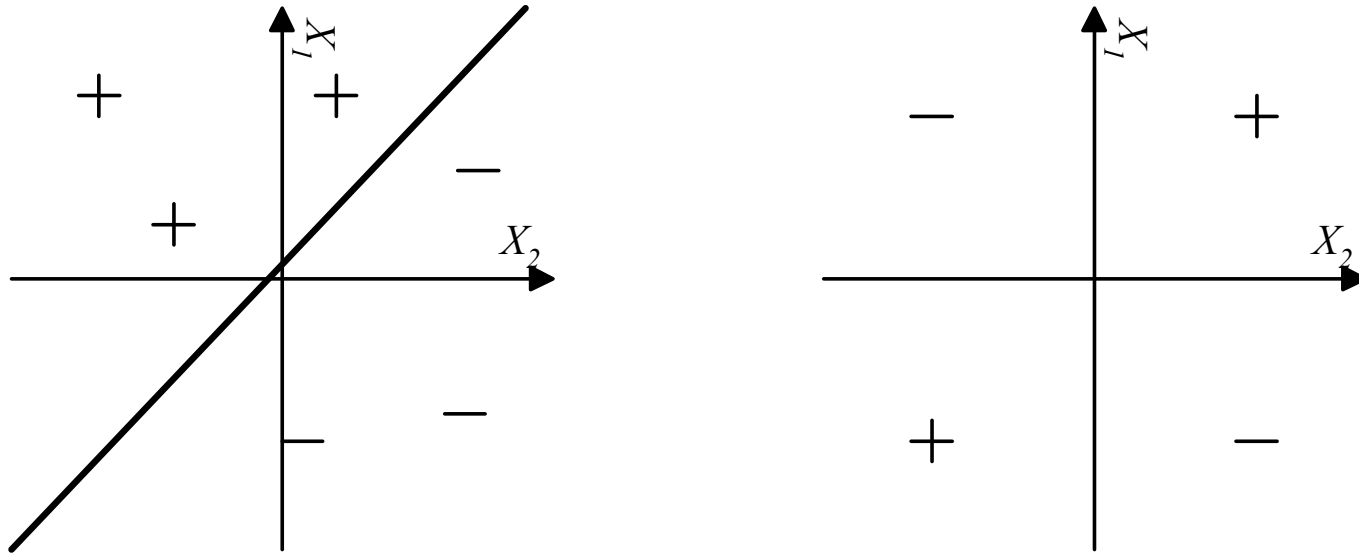
# The Sigmoid Function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



- “rounded” step function
- Unlike step function, can take derivative → helpful for learning

# Perceptron decision boundaries



Represents some useful functions

- What weights represent  $g(x_1, x_2) = AND(x_1, x_2)$ ?

But some functions not representable

- e.g., not linearly separable
- therefore, we will want networks of perceptrons



# How to train your Perceptron

$$w_i \leftarrow w_i + \Delta w_i$$

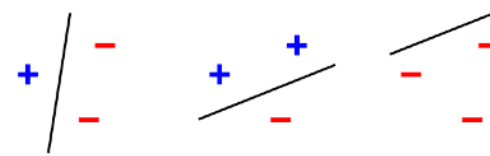
where

$$\Delta w_i = \eta (t - o) x_i$$

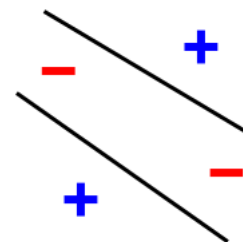
- $t = c(\vec{x})$  is target value
- $o$  is perceptron output
- $\eta$  is small constant (e.g., .1) called learning rate

Can prove it will converge

- If training data is linearly separable
- and  $\eta$  is sufficiently small



linearly separable



not l.s.

# Again: Loss function

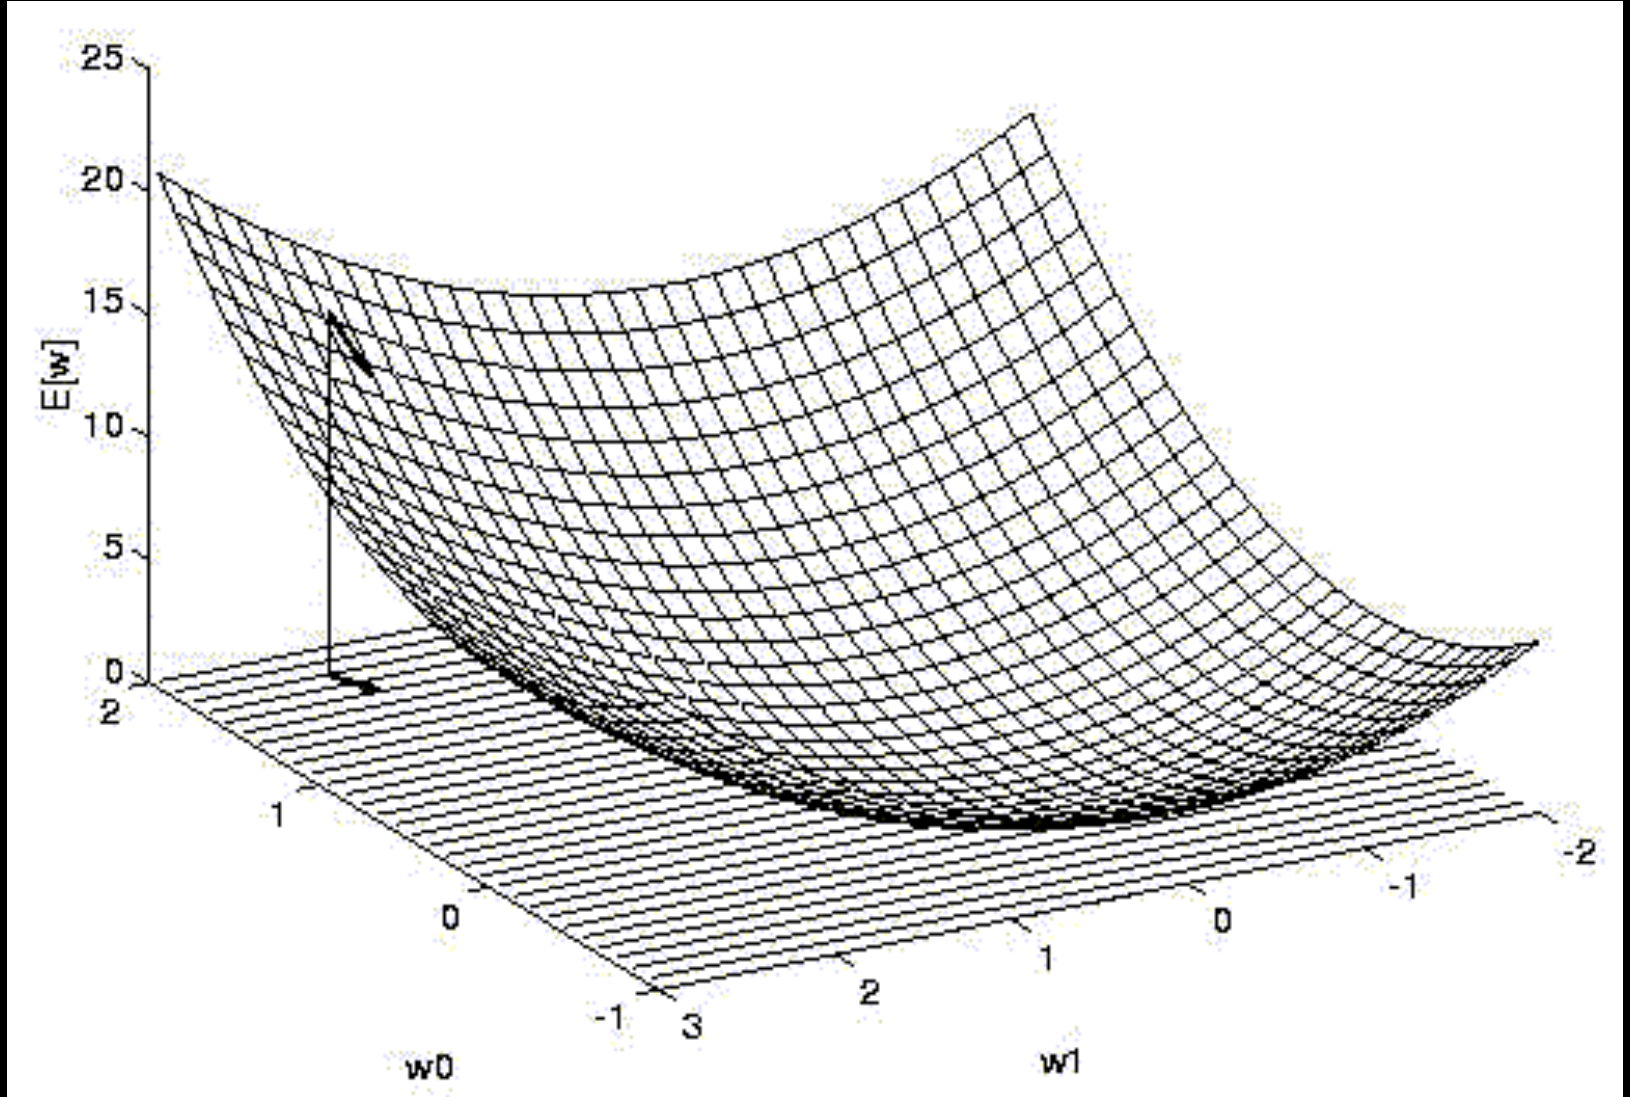
Now we just need to fix the parameters (weights) → Minimize Loss function  $E(\mathbf{w})$ :

$$E(\mathbf{w}) = \sum_i^{events} \left( \underbrace{y_i^{train}}_{\text{true}} - \underbrace{y(x_i)}_{\text{predicted}} \right)^2 \quad \text{i.e. like usual } \chi^2 \text{ in fitting}$$

How to find weights  $w$  that minimize loss function?

- One global fit will usually not work
  - noisy data → many local minima
- Solution: Gradient descent optimization

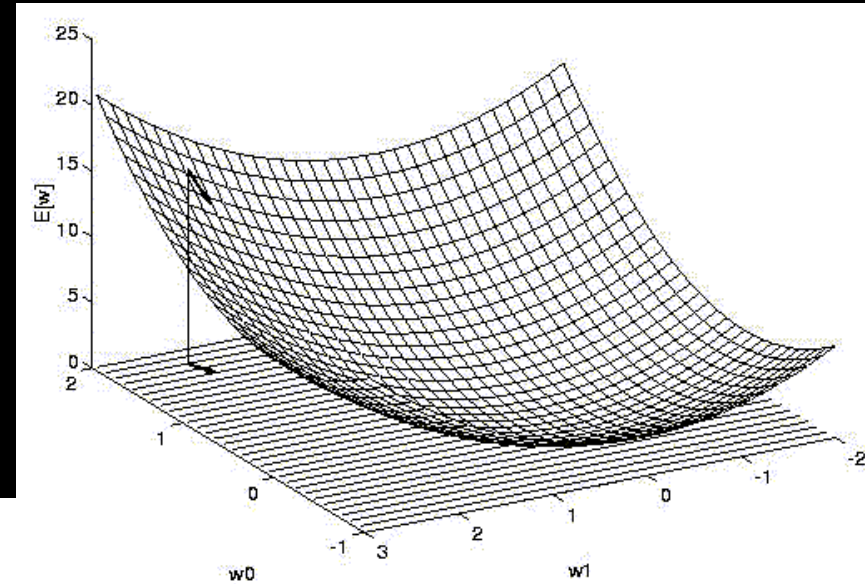
# Gradient Descent



Make steps in the direction of the steepest slope of  $E(\mathbf{w})$

# Gradient Descent

Finding the optimal  
vector of weights  $\mathbf{w}$



Gradient  $\nabla E[\vec{w}] \equiv \left[ \frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$

Slope of  $E(\mathbf{w})$  surface

Training rule:  $\Delta w_i = -\eta \nabla E[\vec{w}]$

small step  $\eta$  down slope  $\nabla E$

i.e., 
$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

# Gradient Descent Training

GRADIENT – DESCENT(*training \_examples*,  $\eta$  )

*Each training examples is a pair of the form  $\langle \vec{x}, t \rangle$  , where  $\vec{x}$  is the vector of input values and  $t$  is the target output value.  $\eta$  is the learning rate (e.g., .05 ).*

- Initialize each  $w_i$  to some small random value
- Until the termination condition is met, do
  - Initialize each  $\Delta w_i$  to zero.
  - For each  $\langle \vec{x}, t \rangle$  in *training \_examples*, do
    - \* Input the instance  $\vec{x}$  and compute output  $o$
    - \* For each linear unit weight  $w_i$ , do
$$\Delta w_i \leftarrow \Delta w_i + \eta (t - o)x_i$$
  - For each linear unit weight  $w_i$ , do
$$w_i \leftarrow w_i + \Delta w_i$$

→ Obtain set of weights for which output is closest to target, as measured by Loss Function

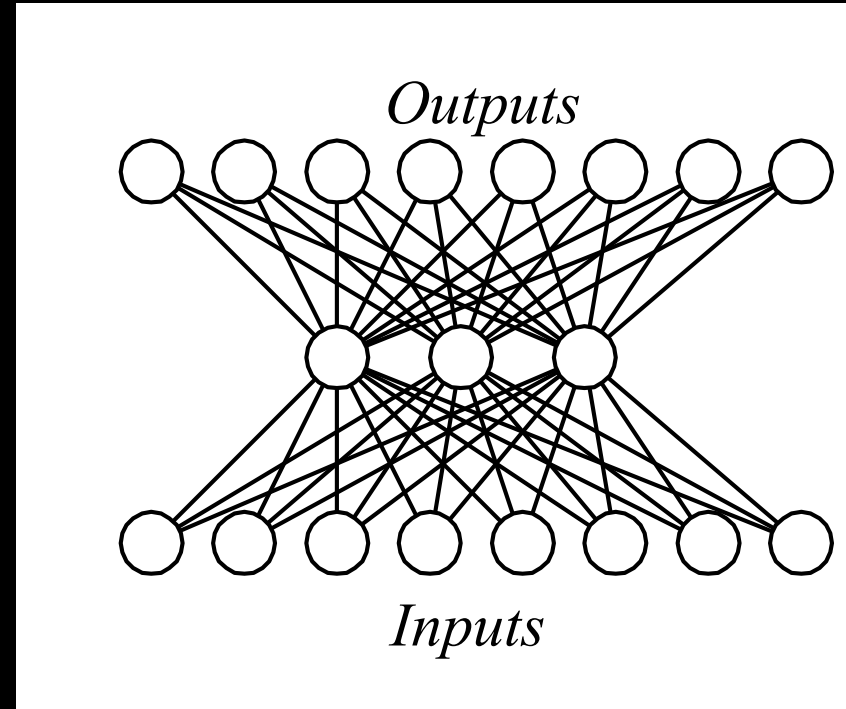
# Gradient Descent

Differentiate using chain rule

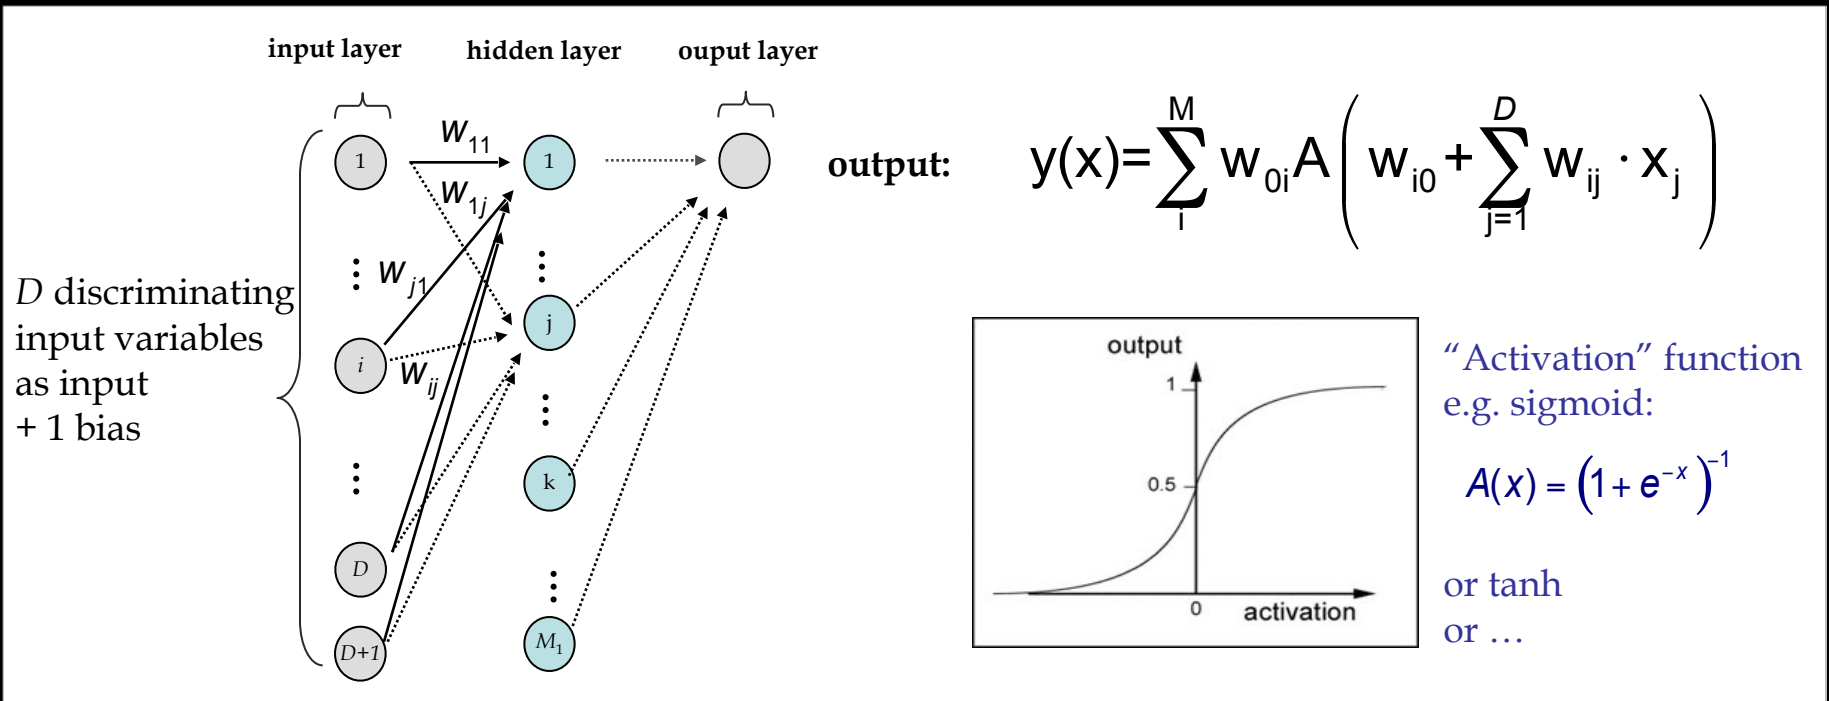
$$\begin{aligned}\frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_d (t_d - o_d)^2 \\&= \frac{1}{2} \sum_d \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\&= \frac{1}{2} \sum_d 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) \\&= \sum_d (t_d - o_d) \frac{\partial}{\partial w_i} (t_d - \vec{w} \cdot \vec{x}_d) \\ \frac{\partial E}{\partial w_i} &= \sum_d (t_d - o_d) (-x_{i,d})\end{aligned}$$

# Networks of perceptrons

- Linear units correspond to *hyperplanes* as decision boundary
- How to approximate arbitrary hyper surfaces?
- → Multi-layer perceptron (MLP)



# Multilayer Perceptron - MLP



- Nodes in hidden layer have “activation functions” whose arguments are linear combinations of input variables → non-linear response to the input
- The output is a linear combination of the output of the activation functions at the internal nodes
- Input to the layers from preceding nodes only → feed forward network (no backward loops)
- It is straightforward to extend this to additional layers



# Backpropagation

- How to change weights for hidden layers?
- Can't take derivative of  $E$  wrt hidden layer weights directly
- Use same idea (gradient descent), but recursively
- Start with output layer, calculate update to weights from hidden layer
- Then update hidden layer weights
  - continue if multiple hidden layers
  - repeat until satisfied with network performance

# Backpropagation

Initialize all weights to small random numbers. Until satisfied, do

- For each training example, do

1. Input the training example and compute the outputs

2. For each output unit  $k$

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$

3. For each hidden unit  $h$

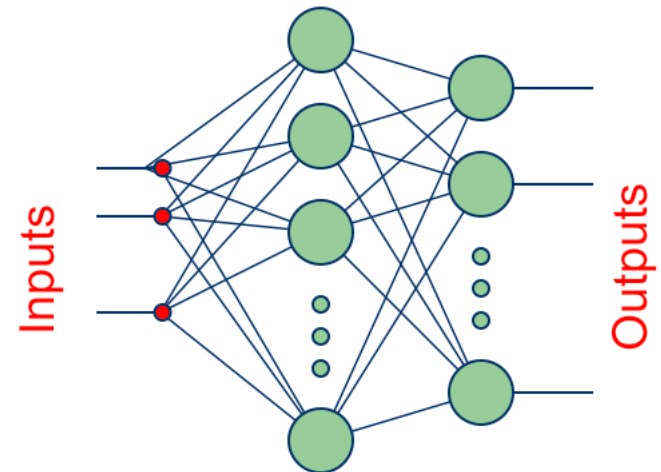
$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{h,k} \delta_k$$

4. Update each network weight  $w_{i,j}$

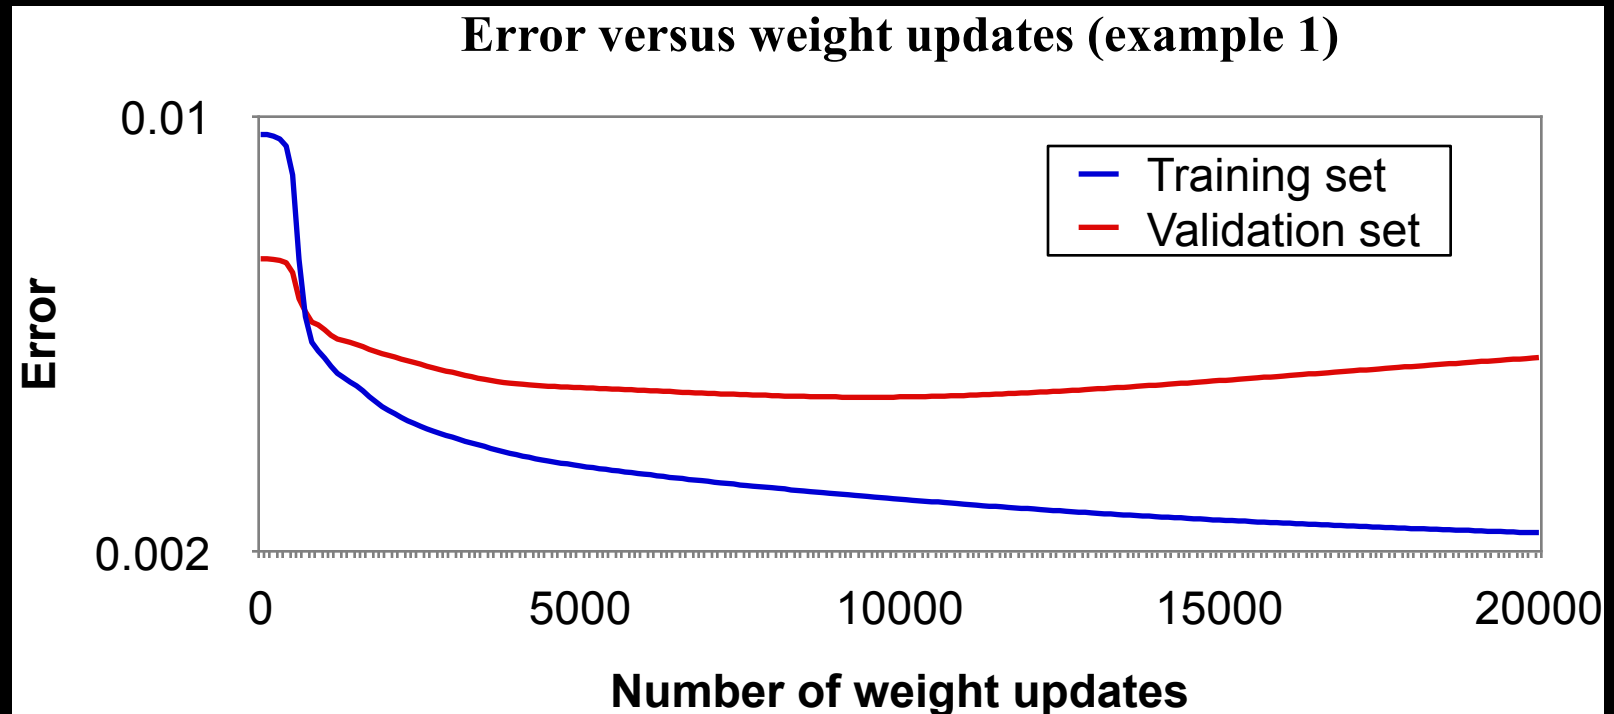
$$w_{i,j} \leftarrow w_{i,j} + \Delta w_{i,j}$$

where

$$\Delta w_{i,j} = \eta \delta_j x_{i,j}$$



# Convergence/Overfitting in ANNs

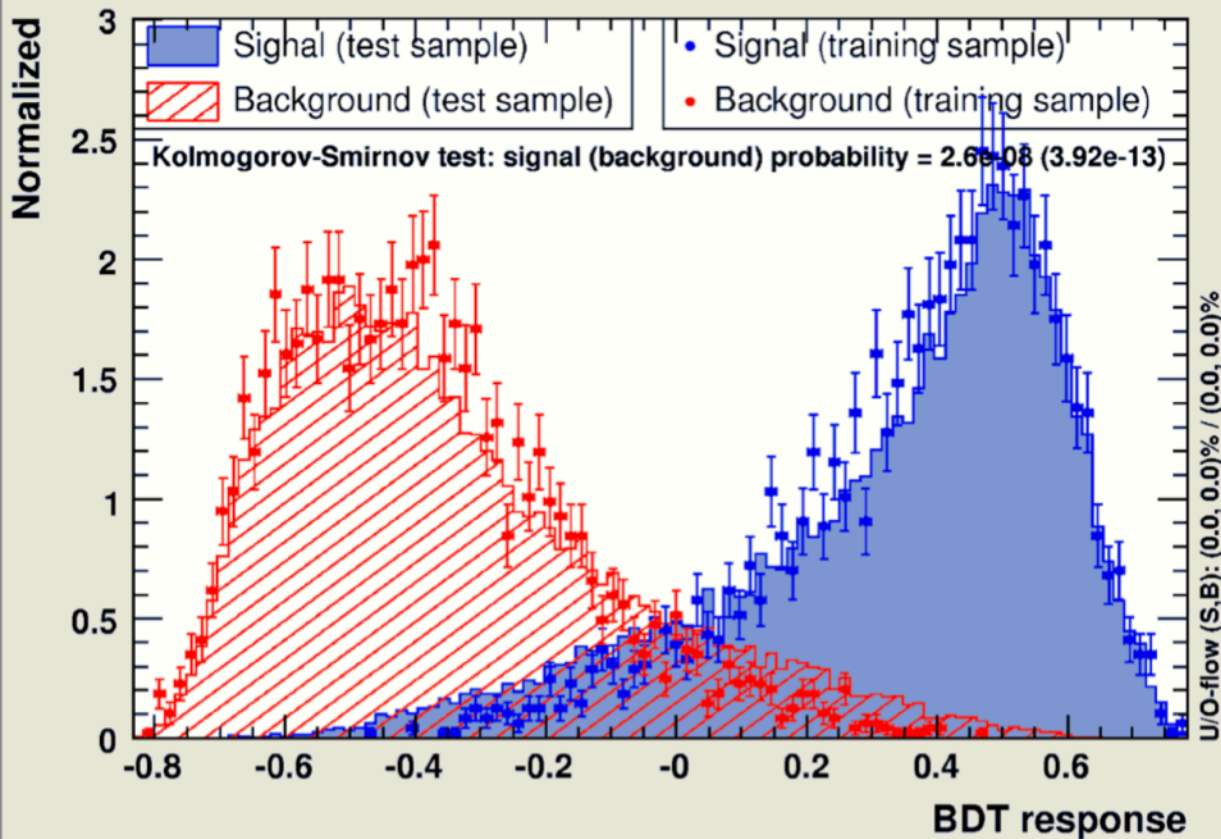


Train on ***Training set***, check results  
on independent ***Validation set (or Test set)***

# Overtraining MVAs

- Check for **overtraining**: classifier output for test *and* training samples

TMVA overtraining check for classifier: BDT



## ■ Remark on **overtraining**

- Occurs when classifier training has too few degrees of freedom because the classifier has too many adjustable parameters for too few training events

- ➔ Sensitivity to overtraining depends on classifier: *e.g.*, **Fisher weak**, **BDT strong**
- ➔ Compare performance between training and test sample to detect overtraining
- ➔ Actively counteract overtraining: *e.g.*, smooth likelihood PDFs, prune decision trees

Let's look at some MLP examples in  
Scikit learn